Enabling HDFS High Availability

1. Go to the HDFS service.
2. Select **Actions** > **Enable High Availability**. A screen showing the hosts that are eligible to run a standby NameNode and the JournalNodes displays.
3. Specify a name for the nameservice or accept the default name **nameservice1** and click **Continue**.
4. In the **NameNode Hosts** field, click **Select a host**. The host selection dialog box displays.
5. In the **JournalNode Hosts** field, click **Select hosts**. The host selection dialog box displays.
6. Check the checkboxes next to an odd number of hosts (a minimum of three) to act as JournalNodes and click **OK**. JournalNodes should be hosted on hosts with similar hardware specification as the NameNodes. Cloudera recommends that you put a JournalNode each on the same hosts as the active and standby NameNodes, and the third JournalNode on similar hardware, such as the JobTracker.
7. Click **Continue**.
8. In the **JournalNode Edits Directory** property, enter a directory location for the JournalNode edits directory into the fields for each JournalNode host. You may enter only one directory for each JournalNode. The paths do not need to be the same on every JournalNode. The directories you specify should be empty, and must have the appropriate permissions.
9. Click **Continue**.
10. Cloudera Manager executes a set of commands that will stop the dependent services, delete, create, and configure roles and directories as appropriate,

create a nameservice and failover controller, and restart the dependent services and deploy the new client configuration.

Configuring Hue to Work with HDFS HA

1. Add the HttpFS role.
    a. Go to the HDFS service.
    b. Click the Instances tab.
    c. Click Add Role Instances.
    d. Click the text box below the HttpFS field. The Select Hosts dialog box displays.
    e. Select the host on which to run the role and click OK.
    f. Click Continue.
    g. Check the checkbox next to the HttpFS role and select Actions for Selected > Start.
2. After the command has completed, go to the **Hue** service.
3. Click the **Configuration** tab.
4. Locate the **HDFS Web Interface Role** property or search for it by typing its name in the Search box.
5. Select the **HttpFS** role you just created instead of the NameNode role, and save your changes.
6. Restart the Hue service.

Upgrading the Hive Metastore to Use HDFS HA

1. Go the Hive service.
2. Select **Actions** > **Stop**.

    Note: You may want to stop the Hue, oozie and Impala services first, if present, as they depend on the Hive service. Click **Stop** to confirm the command.
3. Back up the Hive metastore database.
4. Select **Actions** > **Update Hive Metastore NameNodes** and confirm the command.
5. Select **Actions** > **Start**.
6. Restart the Hue, oozie and Impala services if you stopped them prior to updating the metastore.